Using GCC Auto-Vectorizer

Ira Rosen <ira.rosen@linaro.org>
Michael Hope <michael.hope@linaro.org>

Using GCC Vectorizer

- Vectorization is enabled by the flag -ftree-vectorize and by default at -O3:
 - gcc –O2 –ftree-vectorize myloop.c
 - or gcc –O3 myloop.c
- To enable NEON:
 - -mfpu=neon -mfloat-abi=softfp or -mfloat-abi=hard
- Information on which loops got vectorized, and which didn't and why:
 - -fdump-tree-vect(-details)
 - dumps information into myloop.c.##t.vect
 - -ftree-vectorizer-verbose=[X]
 - dumps to stderr
- More information: http://gcc.gnu.org/projects/tree-ssa/vectorization.html

Other useful flags

- -ffast-math if operating on floats in a reduction computation (to allow the vectorizer to change the order of the computation)
- -funsafe-loop-optimizations if using "unsigned int" loop counters (can be assumed not to overflow)
- -ftree-loop-if-convert-stores more aggressive if-conversion
- --param min-vect-loop-bound=[X] if have loops with a short trip-count
- -fno-vect-loop-version- if worried about code size

What's vectorizable

- Innermost loops
 - countable
 - no control flow

```
for (k = 0; k < m; k ++)

for (j = 0; j < m; j ++)

for (i = 0; i < n; i ++)

a[k][j][i] = b[k][j][i] * c[k][j][i];
```

- independent data accesses
- continuous data accesses

Example of not vectorizable loop:

```
while (a[i] != 8)

if (a[i] != 0)

control flow

a[i] = a[i-1];

b[i+stride] = 0;

access with unknown stride
```

Special features

- vectorization of outer loops
- vectorization of straight-line code
- if-conversion
- multiple data-types and type conversions
- recognition of special idioms (e.g. dot-product, widening operations)
- strided memory accesses
- cost model
- runtime aliasing and alignment tests
- auto-detection of vector size

Examples:

http://gcc.gnu.org/projects/tree-ssa/vectorization.html

GCC Versions

- Current Linaro GCC is based on FSF GCC 4.6
- Once FSF GCC 4.7 is released (in about six months) Linaro GCC will switch to GCC 4.7
- Some of GCC 4.7 vectorizer related features:
 - __builtin_assume_aligned alignment hints
 - vectorization of conditions with mixed types
 - vectorization of bool

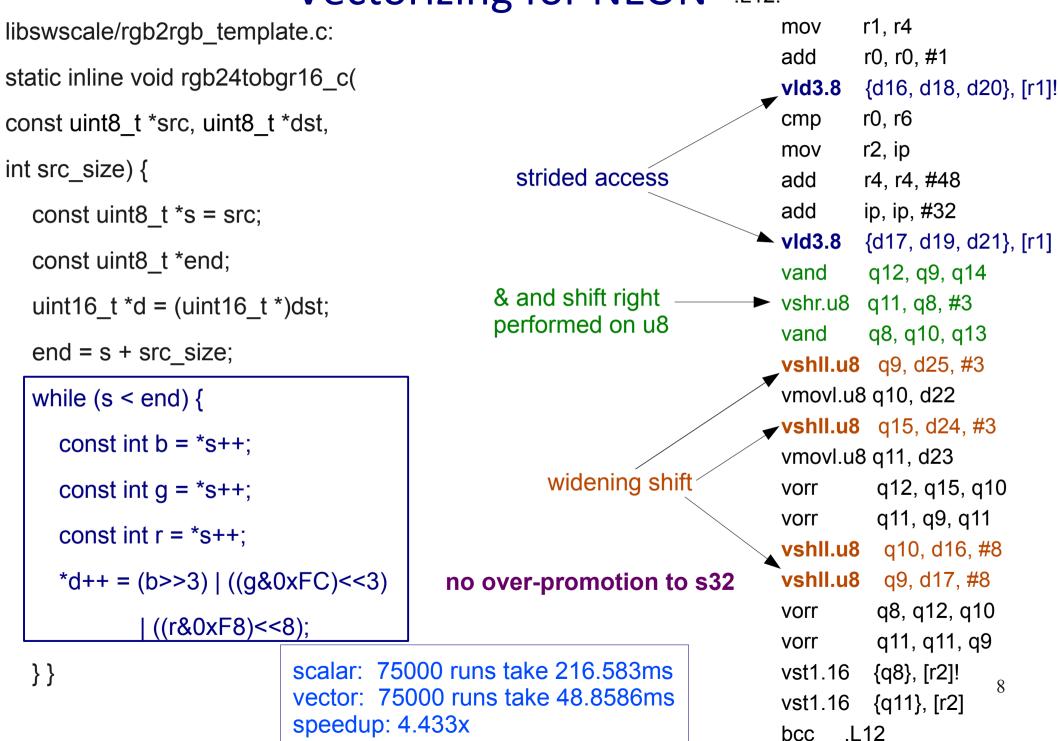
Special features

- vectorization of outer loops
- vectorization of straight-line code
- if-conversion
- multiple data-types and type conversions
- recognition of special idioms (e.g. dot-product, widening operations)
- strided memory accesses
- cost model
- runtime aliasing and alignment tests
- auto-detection of vector size

Examples:

http://gcc.gnu.org/projects/tree-ssa/vectorization.html

Vectorizing for NEON ...



Writing vectorizer-friendly code

- Avoid aliasing problems
 - Use ___restrict__ qualified pointers

```
void foo (int *__restrict__ pInput, int *__restrict__ pOutput)
```

- Don't unroll loops
 - Loop vectorization is more powerful than SLP

```
for (i=0; i<n; i+=4) {
  sum += a[0];
  sum += a[1];
  sum += a[2];
  sum += a[3];
  a += 4;}

for (i=0; i<n; i++)
  sum += a[i];
  sum += a[3];
```

Writing vectorizer-friendly code (cont.)

- Use countable loops, with no side-effects
 - No function-calls in the loop (distribute into a separate loop)

```
for (i=0; i<n; i++)

for (i=0; i<n; i++) {

    if (a[i] == 0) foo();

    if (a[i] == 0) foo();

    b[i] = c[i]; }

    for (i=0; i<n; i++)

    b[i] = c[i];
```

No 'break'/'continue'

```
for (i=0; i<n; i++) {
    if (a[i] == 8) break;
    b[i] = c[i]; }
```

```
for (i=0; i<n; i++)

if (a[i] == 8) {m = i; break;}

for (i=0; i<m; i++)

b[i] = c[i];
```

Writing vectorizer-friendly code (cont.)

- Keep the memory access-pattern simple
 - Don't use indirect accesses, e.g.:

```
for (i=0; i<n; i++)
a[b[i]] = x;
```

Don't use unknown stride, e.g.:

```
for (i=0; i<n; i++)
a[i+stride] = x;
```

- Use "int" iterators rather than "unsigned int" iterators
 - The C standard says that the former cannot overflow, which helps the compiler to determine the trip count.

Some of our recent contributions

- Support of vldN/vstN
- NEON specific patterns: e.g. widening shift
- SLP (straight-line code vectorization) improvements
- RTL improvements:
 - reducing the number of moves and amount of spilling (both for auto- and hand-vectorised code)
 - improving modulo scheduling of NEON code

People

- Linaro Toolchain WG
- Ira Rosen (IRC: irar)

ira.rosen@linaro.org

- auto-vectorizer
- Richard Sandiford (IRC: rsandifo) richard.sandiford@linaro.org
 - NEON back-end/RTL optimizations

Helping us

Send us examples of code that are important to you to vectorize.

Output Example

ex.c:

```
1 #define N 128
2 int a[N], b[N];
3 void foo (void)
4
5
      int i;
6
7
      for (i = 0; i < N; i++)
8
        a[i] = i;
9
10
      for (i = 0; i < N; i+=5)
11
        b[i] = i;
12
```

What's got vectorized:

```
qcc -c -O3 -ftree-vectorizer-verbose=1 ex.c
```

ex.c:7: note: LOOP VECTORIZED. ex.c:3: note: vectorized 1 loops in function.

What's got vectorized and what didn't:

```
gcc -c -O3 -ftree-vectorizer-verbose=2 ex.c
```

ex.c:10: note: not vectorized: complicated access pattern.

ex.c:10: note: not vectorized: complicated access pattern.

ex.c:7: note: LOOP VECTORIZED.

ex.c:3: note: vectorized 1 loops in function.All the details:

gcc -c -O3 -ftree-vectorizer-verbose=9 ex.c or

gcc -c -O3 -fdump-tree-vect-details ex.c